

functionality being added to BioSeq and discusses our hypothesis surrounding performance measures we anticipate being provided by Julia over R. Section 3 discuss the particular algorithms used in the `countOverlaps()` operation as well as discusses the SAMtools API. Section 4 briefly discusses the level of current implementation. Section 5 explores the benchmarking: procedure, data, results, and makes some basic conclusions. Section 6 explains immediate needs before public release. ?

This project seeks to replicate a portion of BioConductor package that is currently available in R for the Julia language. Several key elements of this package have been selected, which includes: basic support for ? genomic data classes (i.e. `IRanges`, `GRanges`, etc.), support to perform union and intersection operations over genes, the ability to import large BAM files containing genomic data, and support for running overlap queries ?. With respect to loading BAM files, we have taken advantage of a Foreign Function Interface (FFI) to take advantage of the SAMtools library.

In order to demonstrate that Julia offers performance benefits over R, we will collect system times on the diverse set of tasks that have been replicated in Julia. In particular, this study will look at the total elapse time of the `countOverlaps()` pipeline as well as the elapse time of each individual set of tasks. Based on the interpreted vs compiled nature of making code comparisons between Julia and R, we make the following hypothesis:

H1. BAM files processing tasks and currently implement data structure operations will show no significant differences in performance times. The use of FFI to call SAMtools is common across both languages and data structure operations are implemented in a naive fashion. Therefore, it is anticipated that the Julia implementation will be faster than R; however, this difference will not be significant.

H2. Performing the `countOverlaps()` operation on the `GRanges` interval object will demonstrate significant performance improvements in Julia that result in improvement in the total elapse time of the pipeline. The `countOverlaps()` operation has been implemented and optimized for running in Julia and is expected to demonstrate performance improvements over R implementation of the same operation.

The purpose of this study is to demonstrate the performance benefits over R that are offered by Julia. These demonstrated benefits will provide support for the continued development of the bioinformatics package `gBioSeq`.

There are no novel algorithms in our implementation, however we do use specific data structures and algorithms to make certain operations efficient. The main example of this is the use of an interval tree. An interval tree can be constructed out of an array of genomic ranges. Then this data structure can be used to efficiently perform queries to find which intervals overlap.

Interval trees are binary trees that contain intervals ?. Each node of the tree contains a center point, an array of intervals that overlap the center point sorted by the start endpoint, and another array of the same intervals sorted by the finish endpoint. The node also stores a pointer to a left node which is a parent of all intervals that are left of the center point. The same is true for a right node and intervals to the right of the center point. A diagram illustrating interval trees can be seen in Figure ?? where the shown intervals are sorted by the start endpoint.

Intervals trees can be constructed by taking the mean of the given intervals as the center point. Then intervals that overlap the center




Fig. 1. An interval tree where the intervals in each node are shown as sorted by the start endpoint. The dashed vertical line indicates the center point of the node.

point are stored in the current node. Intervals to the left or right of the center point are then recursively split off to construct the left and right nodes until no intervals remain. This how process of construction takes $O(n \log(n))$ time.

Interval trees are used to find the intervals that overlap a query interval. This is accomplished by checking if the query interval overlaps with a node's center point. If this is true then all the intervals in that node are found to overlap. If the query interval is left of the center point, the node's forward sorted array of interval is checked if they overlap with the query. Once one of the intervals is found to not overlap the query, the algorithm stops checking more intervals because the intervals are sorted. The same approach is applied to the case when the query interval is right of the center point. After a node is finished, it recursively runs the query against its left or right child nodes as necessary. Performing queries on interval trees takes $O(\log(n))$ time as long as the tree is sufficiently balanced.

2 DISCUSSION

these sections/subsections seem a bit weird to me Feel free to change them... the format is loosely based on the guide; hector suggested doing our best with it, but modify as necessary

2.1 Data

Our data was collected using the system timing routines provided by the standard libraries of Julia and R. Table X shows the average time it takes to read in a test data file and map it to memory, run several `GRange` functions, and complete multiple queries to `countOverlaps()` routine.

5.2 Results

With respect to H1 we found...

With respect to H2 we found...

2.2 Limitations

This implementation developed for this study is very preliminary. The pipeline necessary for performing the `countOverlaps()` operation has been implemented, but rigorous testing is still required to verify the apparent deterministic behaviour beyond the limited testing that has been performed. Additionally, the data that was used to test has been both randomly generated or provided as sample data for a previous project. The purpose of using randomly generated data was to verify the correctness of the data using simple examples and then testing the performance on various sized data sets. Using past data was done to verify correctness of the FFI interface and SAMtools. Correctness and performance is yet to be verified on large real datasets. Finally, considerable time was invested into all three components that were constructed for the `countOverlaps()` pipeline; however, most are not optimized in the same way that `countOverlaps()` was and therefore are sources of potential performance bottlenecks for future users.

2.3 Conclusions

R might not be the best tool for Bioinformatics. Extending the BioSeq package for Julia will allow researchers to utilize an environment that supports the reduction of programming errors through a typing hierarchy, which provides improved performance over R when combined with the JIT compiler provided by Julia. By implementing a portion of BioSeq, we hope to encourage more development of this package by the Bioinformatics community.

TODO: We propose that the bioinformatics community take advantage of, and begin developing packages for, the Julia Language (<http://julialang.org>). Julia provides a strong typing architecture that reduces some of the ambiguity in verifying the correctness of a program. Additionally, this information can be provided to a compiler to make some assurances about the code. The JIT compiler of Julia provides performance benefits over R; however, packages for this language are limited making R superior

in terms of ease of use in - especially for the bioinformatics community. Without support from developers and researchers, this actively developing language will continue to be overlooked by researchers despite its many benefits.

Future Work

This study has demonstrated the preliminary development of expansions to the current BioSeq (Zea, 2013) package being developed for Julia. This study showcases that Julia outperforms R in terms of performance during our selected benchmarking operation. This expansion includes the basic data structures, operations, and interfaces associated with genomic ranges and intervals in R and Bioconductor. However, further work on these objects needs to be performed. Specifically, the genomic ranges are not fully implemented as defined by the R specifications found in the genomic ranges vignette (Carlson et al., 2013) and should be. We aim to complete the functions that perform union and intersection operations across genomic ranges and perhaps a few others as necessary..

We have also demonstrated that a simplified API wrapper for accessing SAMtools, and BAM files, through an FFI can be replicated. The API is very minimal, designed produce GRanges of the data specifically to assist in performing the `countOverlaps()` operation. This API should also be further developed before release **TODO: expand on this**. Finally, general documentation and clean up tasks still need to be performed before a large release of the project can be considered.

ACKNOWLEDGEMENT

The authors would like to thank Dr. Corrada Bravo, of University of Maryland, for his guidance and direction during the course of this project. Additional thanks to the Julia Language and BioSeq communities for providing responsive feedback and assistance with conflict resolution in this early developing environment.